ChatGPT

Lecture 1.2: Opportunity Cost

In economics, every choice requires forgoing some alternative. This fundamental trade-off under scarcity gives rise to **opportunity cost**, the value of the best alternative that must be sacrificed to undertake an action. In precise terms, the *opportunity cost* of an activity is the benefit you would have received from the next-best alternative use of the resources [1]. Equivalently, under constrained choice (scarcity) the **marginal opportunity cost** of using one more unit of a resource in one way is measured by how much of some other good must be given up. Throughout this lecture we build on Lecture 1.1's framework of constrained choice and marginal analysis to formalize and analyze opportunity costs in production, consumption, and intertemporal decisions. We assume agents are rational (with complete, transitive preferences and optimizing behavior) and use both graphical Production Possibility Frontiers (PPFs) and algebraic models to derive results. Where markets are competitive, competitive prices will exactly reflect opportunity costs (first-best); but in second-best settings (e.g. with taxes or externalities) market prices deviate from social opportunity costs. We also sketch applications (consumer trade-offs, firm investment, policy decisions) to illustrate these principles. Canonical references include Varian, Mas-Colell, Samuelson, and the microeconomic literature on equilibrium. Throughout, *opportunity cost* is intimately tied to the slope of constraints: for production the PPF's slope (the **marginal rate of transformation** or MRT), and for consumption the budget line's slope (the price ratio, reflecting the **marginal rate of substitution** or MRS at optimum).

Scarcity and the Definition of Opportunity Cost

By scarcity, resources are limited relative to wants, so choices inevitably involve *trade-offs*. A policy or decision that increases one outcome generally reduces some other. The **opportunity cost** of an action is defined as the value of the best alternative foregone due to that action. In other words:

- **Definition (Opportunity Cost):** *The opportunity cost of any choice is the benefit of the next-best alternative that is given up by making that choice.*

This concept underpins rational decision-making. For example, if a consumer has a limited budget and buys one more unit of good X, the opportunity cost is how much of good Y she must give up. If a firm uses capital in project A, the opportunity cost is the forgone return from project B. Formally, in a constrained optimization (e.g. a Lagrange problem) the Lagrange multiplier on the resource constraint is the *shadow price* or marginal opportunity cost of that resource. Intuitively, opportunity cost is omnipresent – "no free lunch" – because resources (time, money, labor, capital) are scarce.

It is important to note that opportunity cost always reflects *marginal* sacrifice: the cost of the *next* unit of a resource or next alternative. In this marginal analysis, one typically equates marginal benefits to marginal opportunity costs at an optimum. In production, the marginal opportunity cost of expanding output of one good is the amount of another good that must be foregone (captured by the slope of the PPF, or MRT). In consumption, the marginal opportunity cost of buying an extra unit of a good is given by the price ratio (the slope of the budget line) and at optimum equals the marginal rate of substitution. In an intertemporal context, the opportunity cost of consuming now is the foregone future consumption (involving the interest rate or discount factor).

Opportunity cost plays a crucial role in efficiency. In a competitive **General Equilibrium**, the first-order conditions for efficiency equate the marginal rate of substitution (MRS) in consumption to the marginal rate of transformation (MRT) in production for all goods. Equivalently, all agents face the same price ratio, which is the (market-revealed) opportunity cost of goods. When this condition holds, no further Pareto improvements are possible. However, in second-best settings (with distortions), the competitive price ratio may not equal the true marginal social cost, and thus unexploited gains remain.

Throughout this lecture we will illustrate these ideas with formal derivations and graphs. We will cite canonical sources where possible. For example, intermediate theory tells us that on the PPF, the **slope equals the marginal rate of transformation (MRT)**, which "*captures the opportunity cost of producing an extra unit of good 1, in terms of units of good 2 given up*" [2]. Similarly, consumer theory yields $MRS = p_1/p_2$ at optimum [3], meaning the price ratio (the market's opportunity cost of one good) equals the consumer's willingness to trade goods.

In summary, opportunity cost formalizes the trade-off aspect of scarce choice. It tells us *exactly* how much of one thing is sacrificed to get more of another. We now develop this concept in various contexts.

Opportunity Cost in Production: PPF and MRT

Consider an economy (or firm) producing two goods, 1 and 2, using a fixed set of resources (e.g. labor and capital). The **Production Possibility Frontier (PPF)** is the set of output bundles (y_1, y_2) that can be produced with full efficiency. Graphically, the PPF is drawn in (y_1, y_2)-space and typically is concave (bowed outward) from the origin. Under standard assumptions (diminishing marginal products), producing more of good 1 requires increasingly larger reductions in good 2. The concavity reflects *increasing opportunity costs*: to gain one more unit of good 1, we must give up progressively more units of good 2. In fact, "a concave PPF curves outward from the origin, indicating a greater sacrifice of one good as more of the other good is produced" (typically drawn for introductory purposes) [4].

Example 1: Suppose good 1 is produced by labor with diminishing returns: $y_1 = f_1(L_1) = 10\sqrt{L_1}$, and good 2 by another use of labor $y_2 = f_2(L_2) = 6\sqrt{L_2}$. Total labor is $L_1 + L_2 = L$. Inverting these, $L_1 = (y_1/10)^2,\ L_2 = (y_2/6)^2$, and $L = (y_1/10)^2 + (y_2/6)^2 \le \bar{L}$. The efficient frontier is given by the implicit function $L(\, y_1, y_2\,) = \bar{L}$. In parametric form, as we reallocate labor between sectors, the slope (MRT) is

$$MRT = \frac{dy_2}{dy_1}\Big|_{\text{PPF}} = \frac{MP_{L,2}}{MP_{L,1}} \;=\; \frac{\partial f_2/\partial L_2}{\partial f_1/\partial L_1}\ .$$

Here $MP_{L,i}$ are the marginal products. Because $MP_{L,1} = 5/\sqrt{L_1}$ and $MP_{L,2} = 3/\sqrt{L_2}$, one finds $MRT = (3\sqrt{L_1})/(5\sqrt{L_2})$, which when expressed in terms of y_1, y_2 becomes $\frac{9y_1}{25y_2}$. This MRT is increasing as y_1 increases (and y_2 falls), reflecting increasing opportunity cost (Bowley's law of diminishing returns) [5].

More generally, **Definition (Marginal Rate of Transformation):** Along the PPF, the *MRT of good 1 for good 2* is the absolute value of the slope of the frontier:

$$|MRT_{1\to 2}| \;=\; \left|\frac{dy_2}{dy_1}\right|_{\mathrm{PPF}}.$$

Economically, $MRT_{1\to 2}$ measures how many units of good 2 must be sacrificed to produce one additional unit of good 1, holding technology and resource constraints fixed. Thus the MRT *is* the opportunity cost of good 1 in terms of good 2 [1] [2]. For example, if $MRT_{1\to 2} = 4$, then producing one more unit of good 1 requires giving up 4 units of good 2. Conversely, the inverse $|dy_1/dy_2|$ is the opportunity cost of good 2 in terms of good 1.

A classic result ties the MRT to marginal products of inputs. If labor is the only input, reallocating a small amount dL from sector 2 to sector 1 increases y_1 by $MP_{L,1}dL$ and reduces y_2 by $MP_{L,2}dL$. Hence

$$|MRT_{1\to 2}| \;=\; \frac{MP_{L,2}}{MP_{L,1}},$$

the ratio of marginal products, at efficient points. In fact, Varian shows that if the economy is on the PPF and efficient, reallocating any input yields the same ratio. He writes:

> "Suppose we want to reduce production of good 1 and increase production of good 2. One way is to reallocate 1 hour of labor from industry 1 to industry 2. That increases good 2 by MP_2^L and decreases good 1 by MP_1^L. So $\left|dX_2/dX_1\right| = MP_2^L/MP_1^L$. Alternatively reallocating capital yields $\left|dX_2/dX_1\right| = MP_2^K/MP_1^K$. If production is efficient, it must be true that**
>
> $$|MRT| \;=\; \frac{MP_2^L}{MP_1^L} \;=\; \frac{MP_2^K}{MP_1^K},$$
>
> "the same result however inputs are shifted" [6] [7].

Varian then shows in equilibrium that this MRT must equal relative prices p_1/p_2 [8]. In short, at the efficient frontier, the trade-off between goods equals the trade-off of inputs, and competitive prices will align them.

Graphical model: A typical PPF diagram (concave to origin) is shown below. Points on the curve are efficient (all resources used). The slope at each point is the MRT. Moving from one point to another along the frontier trades off outputs.

【56†】 *Figure: A concave Production Possibility Frontier (PPF). The slope of the PPF at any point is the Marginal Rate of Transformation (MRT), the opportunity cost of one good in terms of the other.*

More formally, if the PPF is given by a smooth curve $y_2 = F(y_1)$, then

$$MRT_{1\to 2} \;=\; -\frac{dF}{dy_1}\;,$$

and its absolute value $\left|MRT_{1\to 2}\right|$ equals $\left|dy_2/dy_1\right|$. The negative sign arises because the PPF is downward sloping: to increase y_1 (moving right), y_2 must decrease (moving down). In any case, "the marginal rate of transformation equals the slope of the PPF" [1], capturing precisely the production opportunity cost.

Example: Suppose a firm has 100 units of labor. Good 1 requires 2 units of labor per unit (output $y_1 = L_1/2$), and good 2 requires 1 unit of labor per unit ($y_2 = L_2$). Thus the PPF is $y_1 + (1/2)y_2 = 50$ (since $2y_1 + y_2 = 100$). This is a straight-line (constant-MRT) PPF. The opportunity cost of one more unit of good 1 is $\Delta y_2 = -2\Delta y_1$ (give up 2 of y_2 for one of y_1), so $MRT_{1\to 2} = 2$. Conversely, each extra unit of good 2 costs 0.5 units of good 1. If instead production suffered diminishing returns (concave PPF), the MRT would vary along the curve, typically increasing in magnitude as one good expands. In fact, "making more of one good means making less of another... resources are fully used on the PPF. This tradeoff is measured by the marginal rate of transformation [9]."

Efficiency and Opportunity Cost: In a competitive general equilibrium with production, first-order conditions imply each firm equates marginal rate of technical substitution (MRTS) of inputs to factor price ratios, which in aggregate yields a PPF where the MRT equals the goods price ratio. Simultaneously, each consumer equates MRS to the same price ratio. Varian (as above) summarizes: if firms maximize profit and consumers maximize utility in a price-taking economy, then

$$|MRT| = \frac{p_1}{p_2}, \quad MRS = \frac{p_1}{p_2} \quad \Longrightarrow \quad MRS = MRT.$$

Thus *in equilibrium the consumer's marginal willingness to trade goods (MRS) equals the economy's marginal trade-off in production (MRT)* [10] [1]. This is the formal expression of Pareto efficiency: no reallocation can make someone better off without hurting someone else because every trade-off (opportunity cost) is balanced.

Summarizing production opportunity cost: **the slope of the PPF is the key**. It tells us exactly how output must change across goods. Formally, at a point on the PPF,

$$|MRT_{1\to 2}| \; = \; \frac{MP_2}{MP_1}$$

(where MP_i are marginal products of labor or other input), and it equals the opportunity cost of good 1 (in units of good 2) [6] [2]. A varying (increasing) MRT implies increasing marginal opportunity costs. We will use this notion when analyzing how changes in resources or technology shift the PPF (e.g. richer technology lowers MRTs).

Opportunity Cost in Consumption: Budget Constraint and MRS

Turning to individual choice between two goods, consider a consumer with utility $u(x_1, x_2)$, incomes and prices (p_1, p_2). The consumer faces a budget constraint

$$p_1 x_1 + p_2 x_2 \; = \; I,$$

where I is income. Graphically, the budget line is downward sloping with slope $-p_1/p_2$. Any bundle on the line uses all income; bundles inside are affordable (scarcity binds), those above are unaffordable. The slope $-p_1/p_2$ tells us the *market's* trade-off between goods: giving up p_1/p_2 units of good 2 to get one more unit of good 1 (in order to keep spending constant). Hence p_1/p_2 is the *opportunity cost of good 1 in terms of good 2*.

The consumer maximizes utility subject to the budget. Standard analysis (Lagrange method) yields first-order conditions: if $x_1^*, x_2^* > 0$, then

$$\frac{\partial u/\partial x_1}{\partial u/\partial x_2} \;=\; \frac{p_1}{p_2}.$$

The left side is the **marginal rate of substitution** $MRS_{12} = MU_1/MU_2$, the rate at which the consumer is willing to trade good 2 for good 1 while maintaining equal utility. The right side is the price ratio. At the optimum, the "psychic" trade-off (MRS) equals the "market" trade-off (price ratio) [11] [3]. Intuitively, when $MRS > MRT$ (steeper indifference curve than budget), the consumer would value an extra x_1 more than it costs, so she would buy more x_1; when $MRS < MRT$, she would buy more x_2. Only at equality is she in equilibrium.

Thus **Definition:** The budget constraint's slope $-p_1/p_2$ is the opportunity cost of x_1 in terms of x_2. Equivalently, each extra x_1 requires sacrificing p_1/p_2 units of x_2. At a utility-maximizing bundle, we have

$$MRS_{12} \;=\; \frac{p_1}{p_2},$$

so the consumer's marginal willingness to trade equals the market's required trade-off [3] [11]. In plain language, the consumer spends income until the "bang-for-buck" of each good is equalized: $MU_1/p_1 = MU_2/p_2$. This condition implicitly defines the opportunity cost constraint.

Graphically, one draws indifference curves and a budget line. The point of tangency is the optimal consumption (x_1^*, x_2^*). At that tangency, the slope of the indifference curve (MRS) equals the slope of the budget line (price ratio). That common slope is the *marginal opportunity cost* of one good in terms of the other in the consumer's equilibrium choice.

Example: Let $u(x_1, x_2) = \sqrt{x_1 x_2}$ (Cobb–Douglas) and prices $p_1 = 1,\ p_2 = 2$, income $I = 100$. The FOCs give $MRS = MU_1/MU_2 = (1/2)\sqrt{x_2/x_1}/(1/2)\sqrt{x_1/x_2} = x_2/x_1$. Setting $MRS = p_1/p_2 = 1/2$ yields $x_2/x_1 = 1/2$, so $x_2 = x_1/2$. Substituting into the budget $x_1 + 2x_2 = 100$ gives $x_1 + 2(x_1/2) = 2x_1 = 100$, so $x_1^* = 50,\ x_2^* = 25$. The opportunity cost of one additional unit of good 1 is $p_1/p_2 = 1/2$ units of good 2 (i.e.\ -1 on horizontal, $-1/2$ on vertical move). Indeed, to afford $x_1 = 51$, the consumer must reduce x_2 by 0.5 (to 24.5) to keep spending 100.

Notice that opportunity cost in consumption arises from the budget constraint: for each extra unit of good 1 purchased, income is depleted by p_1, so consumption of good 2 must fall by p_1/p_2 units. In general, at any interior optimum

$$\frac{MU_1}{MU_2} = \frac{p_1}{p_2},$$

so the subjective trade-off (MRS) equals the objective trade-off (price ratio). The price ratio is thus the **marginal opportunity cost** (MRT of consumption) of substituting between goods. We often say *"price equals opportunity cost in competitive markets"*.

To summarize consumption: the key algebraic result is $MRS = p_1/p_2$ at optimum [3]. Geometrically, the **budget line** slope is p_1/p_2, and the consumer's indifference curve slope is MRS. Equating them means the marginal rate of substitution equals the market's marginal rate of transformation (just as in production), reinforcing that opportunity cost ties supply and demand.

Intertemporal Choice: Present vs. Future Consumption

All the above ideas extend naturally to trade-offs over time. In a two-period model, "good 1" can be thought of as consumption in period 1 (today) and "good 2" as consumption in period 2 (tomorrow). Suppose an agent has incomes y_1 (today) and y_2 (tomorrow) and faces an interest rate r for saving/borrowing. The **intertemporal budget constraint** is

$$c_1 + \frac{c_2}{1+r} \ = \ y_1 + \frac{y_2}{1+r},$$

where c_1 is consumption today and c_2 tomorrow in present-value terms. Equivalently, bringing terms together: $(1+r)c_1 + c_2 = (1+r)y_1 + y_2$. In either form, the slope is $-(1+r)$ in c_1-c_2 space. Economically, this slope means **each unit of present consumption costs $(1+r)$ units of future consumption**. As one textbook notes: "the 'price' of present consumption is $1+r$ because spending one dollar today means giving up $(1+r)$ dollars in the future" [12]. Indeed, if you consume one extra unit now ($+\Delta c_1$), then to satisfy the constraint you must decrease future consumption by $\Delta c_2 = -(1+r)\Delta c_1$.

Thus the **opportunity cost of consuming today** is the amount of future consumption forgone at the market interest rate. If $r > 0$, that cost exceeds one: spending $1 now means you lose $1+r$ in future buying power. In present-value language, the budget line can be written $c_1 + c_2/(1+r) = \text{PV income}$, whose slope in c_1-c_2 is $-(1+r)$.

When maximizing intertemporal utility $u(c_1, c_2)$ (often separable with discount factor β), the first-order condition (the Euler equation) equates the marginal rate of substitution between periods to the gross interest rate:

$$\frac{u_1'(c_1)}{u_2'(c_2)} \ = \ 1+r.$$

This is exactly the condition that the subjective rate of impatience (MRS over time) equals the objective trade-off (1+r). In words, at optimum the consumer equates *marginal benefit from saving (future utility)* to *marginal cost of saving (one less unit now yields $1+r$ units later)*.

Example (Intertemporal): Suppose $u(c_1, c_2) = \ln c_1 + \beta \ln c_2$, with $\beta = 1/(1+r)$, incomes y_1, y_2. The Euler condition $\frac{1}{c_1} = \beta \frac{1}{c_2}(1+r)$ yields $c_2 = \beta(1+r)c_1 = c_1$. Combined with the budget $c_1 + c_2/(1+r) = y_1 + y_2/(1+r)$, one solves $c_1 = c_2 = (y_1 + y_2/(1+r))/(1 + 1/(1+r))$. The opportunity cost interpretation: if $r = 0.1$, then to consume one extra $ of c_1, you must reduce c_2 by $1.1. The consumer smooths consumption so that marginal utilities are equalized across time, taking into account the interest rate as the exchange rate between periods.

In summary, the intertemporal trade-off is analogous to the goods trade-off: **today's consumption vs tomorrow's**. The budget line has slope $-(1+r)$, so $1+r$ *is the opportunity cost (MRT) of current consumption in terms of future consumption*. When the consumer optimizes, $MRS_{1,2} = 1+r$, equating time-preference with the market trade-off. This is the cornerstone of consumption-saving theory. If r

changes (e.g. monetary policy), the opportunity cost changes, pivoting the budget line around the endowment point [13] [14].

Opportunity Cost in Investment and Capital Allocation

The concept of opportunity cost of capital is closely related. A firm or planner deciding how much to invest vs consume (or between alternative projects) weighs the marginal benefit of investment against the interest rate (the return on the next-best investment). If the real interest rate is r, then investing \$1 today (instead of consuming) will yield \$(1+r) tomorrow. Therefore the opportunity cost of using \$1 of the firm's funds today is \$(1+r) worth of output in the future.

In corporate finance, this appears as the *hurdle rate*: a company will only invest in a project if its expected return exceeds the opportunity cost of capital (the best alternative return). Similarly, a government allocating budget between projects (e.g. infrastructure vs education) faces an implicit trade-off: spending an extra dollar on one program means losing that dollar's future value on another program. In dynamic models (e.g. the neoclassical growth model), the user cost of capital equals $r + \delta$, the marginal product of capital must equal this cost in first-best equilibrium. Thus capital allocation decisions obey the same principle: invest up to the point where $MP_K = r + \delta$, equating marginal benefit and opportunity cost.

Example (Firm Investment): A firm has earnings E and can either pay them as dividends or reinvest. If the interest rate (and market return) is 5%, then foregoing \$1 of dividend (investing it instead) costs the firm \$1 in present consumption and yields 1.05 next period. The opportunity cost of investing now is the forgone \$1 (plus 5% of missed immediate return). The firm will invest until its marginal productivity of capital equals 5%.

Thus, opportunity cost of capital **is essentially the market rate of return**. When making capital allocation (across time or sectors), rational agents use the interest rate as the exchange rate. Formally, one can include capital stock K in the PPF analysis (as in a two-good model of consumption vs capital goods), leading to a three-dimensional PPF. The marginal trade-off between current and future output (the marginal transformation rate in the time dimension) again involves the interest rate.

General Equilibrium and Economy-Wide Trade-Offs

So far we have looked at single-decision problems. In a **general equilibrium** setting with many agents and goods, opportunity costs emerge as economy-wide trade-offs. Suppose an economy produces many goods with finite factors. One can aggregate technology into an **aggregate PPF** that represents all possible output combinations. If factors are mobile and markets clear, the economy's achievable output frontier is convex and downward-sloping in multi-dimensional space.

In general equilibrium, Walras's law and efficiency imply that all such marginal trade-offs coincide. Concretely, consider two sectors (1 and 2) producing goods 1 and 2 with the same factors. Reallocating labor from sector 2 to 1 will raise y_1 and lower y_2 by certain amounts. The ratio $|dy_2/dy_1|$ is the economy's MRT at that point (as in the previous section). Any such movement affects *aggregate* outputs. Because all inputs market-clearing, the economy is on its PPF. Similarly, in consumption equilibrium each consumer's MRS aligns with the common price ratio.

The key outcome is that a competitive equilibrium with production satisfies the **First Welfare Theorem**: it is Pareto-efficient. In terms of opportunity cost, this means no reallocation of consumption or production can improve someone without hurting another, because $MRS = MRT$ for every good pair [10] [1]. Thus the opportunity cost of producing one good in terms of another (the economy's MRT) equals the opportunity cost that consumers are willing to pay (their MRS) – and both are captured by the relative prices.

From a multi-sector perspective, opportunity cost also appears in trade and specialization. In an open economy, each country's PPF determines its comparative advantage. A country specializes in the good with lower opportunity cost (flatter PPF slope) and trades at the world price. The price ratio acts as a universal opportunity cost. The **Edgeworth box** of production can illustrate two-country or two-good economies: every point on the combined PPF corresponds to a distribution of resources, and the common tangent yields trade volumes.

Foregone output across sectors: When analyzing a policy that shifts resources (e.g. labor tax affecting manufacturing vs services), one can quantify the lost output in each sector. For instance, if sector A uses one more worker, sector B loses MP_B output; thus the marginal social cost of additional labor in A is the product lost in B. This is exactly the PPF trade-off. In some models one can even draw an economy-wide PPF (such as in a two-good Heckscher–Ohlin model) and measure MRT graphically.

In short, **general equilibrium prices equal opportunity costs**. The price of a resource or good in one use reflects its marginal value in all uses. This uniformity underlies factor-price equalization and common slopes. Mas-Colell, Whinston, and Green (1989) note that Pareto efficiency requires (i) equal MRS across consumers and (ii) equal MRS to MRT for goods [10] [1]. Debreu's theory of value (1959) shows that equilibrium prices (shadow values) support this efficient allocation. Thus in a fully functioning economy, opportunity costs are embodied in market prices, which in turn coordinate all sectors.

First-Best versus Second-Best Opportunity Costs

In the ideal "first-best" world of perfect competition and no distortions, **market prices are exact opportunity costs**. That is, the price of a good or input reflects the true marginal social cost of using it. The First Welfare Theorem assures that optimizing agents achieve $MRS = MRT = p$, so no mismatches. However, real economies often have *distortions*: taxes, subsidies, monopolies, externalities, or missing markets. In such "second-best" settings, the privately faced trade-off differs from the social one.

For example, suppose good 1 is taxed at rate t, so consumers pay $(1 + t)p_1$. The consumer equates $MRS_{12} = (1 + t)p_1/p_2$, whereas firms still face output price p_1. Now $MRS \neq MRT$ because $p_1 \neq (1 + t)p_1$. The consumer's opportunity cost of good 1 (in terms of good 2) is $(1 + t)p_1/p_2$, but the producer's opportunity cost is p_1/p_2. The gap $(1 + t)$ reflects deadweight loss. As a result, the equilibrium is inefficient (second-best). In effect, the **shadow price** of consumption (the true social opportunity cost) is lower than the market price paid by consumers.

In another example, an externality (say pollution from producing good 1) means the social cost of producing y_1 exceeds the private cost. The private firm's PPF is too optimistic, understating the true trade-off. Correcting the externality (Pigou tax) would align private and social MRTs. In these cases, one speaks of *shadow prices* or *implicit costs* that reflect social opportunity costs.

Formally, in a constrained planner problem (maximize welfare subject to distortion constraints), the Lagrange multipliers on constraints give “shadow prices” of goods – the true social opportunity costs. The **Second Welfare Theorem** (with constraints) does not hold in general, but one can often compute the *second-best optimum* by weighting goods appropriately. Without going into technical proofs, the key message is: **if one market fails, opportunity costs are no longer equalized by prices**. Agents make decisions based on the distorted price, not the true cost to others, so adjustments are needed (taxes, subsidies, regulation) to restore the real trade-offs.

First vs. second best in policy: A classic description: In first best, the opportunity cost of public funds is 1 (one more dollar of spending costs one dollar of output). In second best (with e.g. pre-existing taxes), the marginal social cost of raising a dollar may exceed one. Thus, a dollar spent on program A has an opportunity cost of not funding B *plus* the inefficiency from raising taxes to pay for A.

In terms of allocation, **first-best opportunity costs** are those under pure efficiency: e.g. the derivative of a social welfare function or production function. **Second-best opportunity costs** include adjustment terms. For example, in production under a distortion, the marginal benefit of shifting factors is different. One can think of second-best MRT: the *social* MRT (reflecting all externalities) vs the *private* MRT (just technology).

In summary, first-best equals market opportunity cost. Second-best requires adjusting for distortions. If we define opportunity cost as *value of the foregone alternative*, then in second-best situations the market price no longer gives that value. The difference must be measured by shadow pricing. (Mas-Colell et al. devote much discussion to shadow prices in constrained optimization and general equilibrium. See also Samuelson’s treatment of Kaldor–Hicks and compensation, which implicitly uses opportunity cost logic in welfare comparisons.)

Applied Examples

To illustrate these concepts, we briefly discuss a few examples:

- **Consumer Trade-Off (Goods):** A student has \$50 per week to spend on pizza (\$1 each) and movies (\$2 each). The opportunity cost of one movie ticket is 2 pizzas (since \$2 could buy 2 pizzas). If her utility is $U(p,m) = \sqrt{p\,m}$, the FOCs give $m/p = p_P/p_M = 2$. Solving with the budget constraint yields her optimal bundle (e.g. 10 pizzas, 5 movies). Here $MRS_{pm} = p/m = 2$ matches the price ratio $1/2$ inverted (since pizzas cost less), consistent with $MRS = p_P/p_M$.

- **Labor–Leisure Trade-Off:** An individual has 24 hours to allocate between leisure and work. If the wage is w, then the opportunity cost of one hour of leisure is w units of consumption (foregone wage). If consumption goods and leisure are the two “goods” in our model, then at optimum $MRS_{l,c} = w$. Thus higher wages raise the opportunity cost of leisure, shifting the labor supply curve (as in Varian’s statement that *“the wage rate is the opportunity cost of leisure”*).

- **Firm Investment Choice:** A firm can either invest \$10,000 now or use it for marketing. If the firm’s own projects yield 8% and the market rate is 5%, the opportunity cost of investing in, say, a low-return project is the higher 8% forgone opportunity. In practice, firms compare internal return rates to a hurdle (often set near market cost of capital). The marginal opportunity cost of \$1 invested is the forgone return from the next best project.

- **Policy Trade-Off (Budget Allocation):** A government has a fixed budget to allocate between infrastructure and education. Suppose spending \$1M on infrastructure yields 0.5 units of output per year, while \$1M on education yields 0.4 units. The PPF between infrastructure and education spending is downward sloping. If initially spending is at an interior point, the government should equate the marginal benefit ratios: it will shift spending from the lower-return program to the higher-return one until 0.5 (for infra) equals 0.4 (for edu), adjusted for shadow values. The opportunity cost of reallocating budget is the lost output in the program cut (the slope of the spending-PPF).

- **Intertemporal Example:** A retiree can either spend down savings today or save for more consumption in retirement. If the interest rate is 3%, then each \$1 of extra consumption today costs \$1.03 of future consumption. The retiree's MRS between future and present consumption will be set equal to 1.03 at optimum. We can solve a two-period problem with given utility $\ln c_0 + 0.97 \ln c_1$ and endowments to find the exact split.

These examples all hinge on opportunity cost: every extra unit of choice costs some other benefit. In lecture exercises, one typically sets up the optimization, writes the Lagrangian, and derives the first-order condition equating marginal trade-offs to opportunity costs.

Conclusion

Opportunity cost is a unifying concept that quantifies *trade-offs*. In production it is captured by the slope of the PPF (the MRT), in consumption by the slope of the budget line (the price ratio), and in intertemporal choice by the interest rate (the present-future trade-off). Under rationality and competitive markets, optimizers equate marginal benefit to opportunity cost at the margin. This yields conditions like $MRS = MRT$ (first-best). In equilibrium, market prices relay opportunity costs across sectors (Walras's law of one price). In settings with distortions, one must distinguish private opportunity cost (market price) from social opportunity cost (shadow price).

Throughout economic analysis, asking "What must I give up to get one more of this?" is central. As one source puts it: *"The Marginal Rate of Transformation (MRT)... indicates how many units of one good you must sacrifice to produce an additional unit of another while keeping resources... constant"* [1]. In every model of choice under scarcity, understanding and computing these marginal opportunity costs is key to predicting behavior and evaluating efficiency.

Sources: Standard intermediate micro and general equilibrium texts discuss opportunity cost in exactly these terms (e.g. Varian (1992), Mas-Colell, Whinston & Green (1995)). Varian explicitly defines MRT as the PPF slope and shows $MRT = MP_2/MP_1$ [6] [7], and that in equilibrium $MRT = MRS = p_1/p_2$ [10]. Consumer theory similarly shows $MRS = p_1/p_2$ [3], tying the price ratio to opportunity cost. We have also drawn on pedagogical expositions (e.g. Makler's EconGraphs [2] [12]) and Investopedia [1] [15] for clear definitions of MRT and opportunity cost.

---

[1] [9] [15] Understanding Marginal Rate of Transformation (MRT) and Its Economic Impact
https://www.investopedia.com/terms/m/marginal_rate_transformation.asp

[2] [5] The Marginal Rate of Transformation: the Slope of the PPF - EconGraphs
https://www.econgraphs.org/textbooks/intermediate_micro/scarcity_and_choice/ppf/mrt

[3] consumertheory_2004.dvi
https://web.stanford.edu/~jdlevin/Econ%20202/Consumer%20Theory.pdf

[4] Lecture 3: The 2x2x2 Heckscher-Ohlin-Samuelson Model
https://isabellemejean.com/lecture%203.pdf

[6] [7] [8] [10] yorku.ca
http://www.yorku.ca/bucovets/2350/lectures/33.pdf

[11] 14.03/14.003 Fall 2016 Lecture 4 Notes
https://ocw.mit.edu/courses/14-03-microeconomic-theory-and-public-policy-fall-2016/662896910b5530e160224afe6ac30752_MIT14_03F16_lec4.pdf

[12] [13] [14] The Intertemporal Budget Constraint - EconGraphs
https://www.econgraphs.org/textbooks/econ51winter25/unit1/lecture3/budget_constraint